

# LESSON NAME

## Intro

### Prerequisites

- FIXME
- ...
- ...

20 min filename

## Episode template

### Questions

- What syntax is used to make a lesson?
- How do you structure a lesson effectively for teaching?
- `questions` are at the top of a lesson and provide a starting point for what you might learn. It is usually a bulleted list.

### Objectives

- Show a complete lesson page with all of the most common structures.
- ...

This is also a holdover from the carpentries-style. It could usually be left off.

The introduction should be a high level overview of what is on the page and why it is interesting.

The lines below (only in the source) will set the default highlighting language for the entire page.

## Section

A section.

### Discussion

Skip to content

g.

- Another discussion topic

# Section

```
print("hello world")
# This uses the default highlighting language
```

```
print("hello world")
```

## Exercises: description

|  Exercise Topic-1: imperative description of exercise

Exercise text here.

|  Solution

Solution text here

## Summary

A Summary of what you learned and why it might be useful. Maybe a hint of what comes next.

## See also

- Other relevant links
- Other link

|  Keypoints

- What the learner should take away
- point 2
- ...

This is another holdover from the carpentries style. This perhaps is better done in a "summary" section.

# Cython

Cython is a superset of Python that additionally supports calling C functions and declaring C types on variables and class attributes. It is also a versatile, general purpose compiler. Since it is a superset of Python syntax, nearly all Python code, including 3rd party Python

[Skip to content](#) → valid Cython code. Under Cython, source code gets translated into optimized compiled as Python extension modules.

Developers can either:

- prototype and develop Python code in IPython/Jupyter using the `%%cython` magic command (**easy**), or
- run the `cython` command-line utility to produce a `.c` file from a `.py` or `.pyx` file, which in turn needs to be compiled with a C compiler to an `.so` library, which can then be directly imported in a Python program (**intermediate**), or
- use `setuptools` or `meson` with `meson-python` to automate the aforementioned build process (**advanced**).

Herein, we restrict the discussion to the Jupyter-way of using the `%%cython` magic. A full overview of Cython capabilities refers to the [documentation](#).

### Important

Due to a [known issue](#) with `%%cython -a` in `jupyter-lab` we have to use the `jupyter-nbclassic` interface for this episode.

## Python: Baseline (step 0)

### Demo: Cython

Consider a problem to integrate a function:

$$\int_a^b (x^2 - x) dx$$

which can be numerically approximated as the following sum:

$$\int_a^b \approx \Delta x \sum_{i=0}^{N-1} (x_i^2 - x_i)$$

where  $a \leq x_i \leq b$ , and all  $x_i$  are uniformly spaced apart by  $\Delta x = (b - a) / N$ .

**Objective:** Repeatedly compute the approximate integral for 1000 different combinations of  $(a)$ ,  $(b)$  and  $(N)$ .

Python code is provided below:

[Skip to content](#)

```

import numpy as np

def f(x):
    return x ** 2 - x

def integrate_f(a, b, N):
    s = 0
    dx = (b - a) / N
    for i in range(N):
        s += f(a + i * dx)
    return s * dx

def apply_integrate_f(col_a, col_b, col_N):
    n = len(col_N)
    res = np.empty(n, dtype=np.float64)
    for i in range(n):
        res[i] = integrate_f(col_a[i], col_b[i], col_N[i])
    return res

```

We generate a dataframe and apply the `apply_integrate_f()` function on its columns, timing the execution:

```

import pandas as pd

df = pd.DataFrame(
{
    "a": np.random.randn(1000),
    "b": np.random.randn(1000),
    "N": np.random.randint(low=100, high=1000, size=1000)
})

%timeit apply_integrate_f(df['a'], df['b'], df['N'])
# 101 ms ± 736 µs per loop (mean ± std. dev. of 7 runs, 10 loops each)

```

## Cython: Benchmarking (step 1)

In order to use Cython, we need to import the Cython extension:

```
%load_ext cython
```

As a first cythonization step, we add the cython magic command (`%%cython -a`) on top of Jupyter notebook by simply compiling the Python code using Cython without any changes.

[Skip to content](#)

[n below:](#)

```

%%cython -a

import numpy as np

def f_cython_step1(x):
    return x * (x - 1)

def integrate_f_cython_step1(a, b, N):
    s = 0
    dx = (b - a) / N
    for i in range(N):
        s += f_cython_step1(a + i * dx)
    return s * dx

def apply_integrate_f_cython_step1(col_a, col_b, col_N):
    n = len(col_N)
    res = np.empty(n, dtype=np.float64)
    for i in range(n):
        res[i] = integrate_f_cython_step1(col_a[i], col_b[i], col_N[i])
    return res

```

Yellow lines hint at Python interaction.

Click on a line that starts with a "+" to see the C code that Cython generated for it.

```

01:
+02: import numpy as np
03:
+04: def f_cython_step1(x):
+05:     return x * (x - 1)
06:
+07: def integrate_f_cython_step1(a, b, N):
+08:     s = 0
+09:     dx = (b - a) / N
+10:    for i in range(N):
+11:        s += f_cython_step1(a + i * dx)
+12:    return s * dx
13:
+14: def apply_integrate_f_cython_step1(col_a, col_b, col_N):
+15:     n = len(col_N)
+16:     res = np.empty(n, dtype=np.float64)
+17:     for i in range(n):
+18:         res[i] = integrate_f_cython_step1(col_a[i], col_b[i], col_N[i])
+19:     return res

```

ANNOTATED CYTHON CODE OBTAINED BY RUNNING THE CODE ABOVE. THE YELLOW COLORING IN THE OUTPUT SHOWS US THE AMOUNT OF PURE PYTHON CODE.

Our task is to remove as much yellow as possible by *static typing*, i.e. explicitly declaring arguments, parameters, variables and functions.

[Skip to content](#) The Python code just using Cython, and it may give either similar or a slight increase in performance.

```
%timeit apply_integrate_f_cython_step1(df['a'], df['b'], df['N'])
# 102 ms ± 2.06 ms per loop (mean ± std. dev. of 7 runs, 10 loops each)
```

## Cython: Adding data type annotation to input variables (step 2)

Now we can start adding data type annotation to the input variables as highlighted in the code example/cython below:

Pure Python

Cython

```
%%cython -a

import cython
import numpy as np

def f_cython_step2(x: cython.double):
    return x ** 2 - x

def integrate_f_cython_step2(a: cython.double, b: cython.double, N: cython.long):
    s = 0
    dx = (b - a) / N
    for i in range(N):
        s += f_cython_step2(a + i * dx)
    return s * dx

def apply_integrate_f_cython_step2(
    col_a: cython.double[:],
    col_b: cython.double[:],
    col_N: cython.long[:],
):
    n = len(col_N)
    res = np.empty(n, dtype=np.float64)
    for i in range(n):
        res[i] = integrate_f_cython_step2(col_a[i], col_b[i], col_N[i])
    return res
```

[Skip to content](#)

```
# this will not work
#%timeit apply_integrate_f_cython_step2(df['a'], df['b'], df['N'])

# this command works (see the description below)
%timeit apply_integrate_f_cython_step2(df['a'].to_numpy(), df['b'].to_numpy(), df['N'])
# 34.3 ms ± 537 µs per loop (mean ± std. dev. of 7 runs, 10 loops each)
```

### Warning

You can not pass a Series directly since Cython definition is specific to an array. Instead we should use `Series.to_numpy()` to get the underlying NumPy array which works nicely with Cython.

### Note

Cython uses the normal C syntax for types and provides all standard ones, including pointers. Here is a list of some primitive C data types (refer to Cython's documentation on [Types](#)):

Cython type identifier	Pure Python dtype
<code>char</code>	<code>cython.char</code>
<code>int</code>	<code>cython.int</code>
<code>unsigned int</code>	<code>cython.uint</code>
<code>long</code>	<code>cython.long</code>
<code>float</code>	<code>cython.float</code>
<code>double</code>	<code>cython.double</code>
<code>double complex</code>	<code>cython.doublecomplex</code>
<code>size_t</code>	<code>cython.size_t</code>

Using these data types, we can also annotate arrays (see [Typed Memoryviews](#)):

- 1D `np.float64` array would be equivalent to `cython.double[:]`,
- 2D `np.float64` array would be equivalent to `cython.double[:, :]` and so on...

### Important

to quote the [Cython documentation](#),

#### **Typing is not a necessity**

Providing static typing to parameters and variables is convenience to speed up your code, but it is not a  
necessity where and when needed. In fact, typing can slow down your code in the case where the

[Skip to content](#) allow optimizations but where Cython still needs to check that the type of some object  
is consistent with the declared type.

# Cython: Adding data type annotation to functions (step 3)

Next step, we further add type annotation to functions. There are three ways of declaring functions:

- `def` - Python style:
  - Called by Python or Cython code, and both input/output are Python objects.
  - Declaring argument types and local types (thus return values) can allow Cython to generate optimized code which speeds up the execution.
  - Once types are declared, a `TypeError` will be raised if the function is passed with the wrong types.
- `@cython.cfunc` or `cdef` - C style:
  - `cdef` functions are called from Cython and C, but not from Python code.
  - Cython treats functions as pure C functions, which can take any type of arguments, including non-Python types, e.g., pointers.
  - This usually gives the *best performance*.
  - However, one should really take care of the functions declared by `cdef` as these functions are actually writing in C.
- `@cython.ccall` or `cpdef` - C/Python mixed style:
  - `cpdef` function combines both `cdef` and `def`.
  - Cython will generate a `cdef` function for C types and a `def` function for Python types.
  - In terms of performance, `cpdef` functions may be as *fast as* those using `cdef` and might be as slow as `def` declared functions.

Pure Python

Cython

```

%%cython -a

import cython
import numpy as np


@cython.cfunc
def f_cython_step3(x: cython.double):
    return x ** 2 - x

@cython.cfunc
def integrate_f_cython_step3(a: float, b: float, N: int):
    s = 0
    dx = (b - a) / N

    for i in range(N):
        s += f_cython_step3(a + i * dx)
    return s * dx

@cython.ccall
def apply_integrate_f_cython_step3(
    col_a: cython.double[:],
    col_b: cython.double[:],
    col_N: cython.long[:]
):
    n = len(col_N)
    res = np.empty(n, dtype=np.float64)
    for i in range(n):
        res[i] = integrate_f_cython_step3(col_a[i], col_b[i], col_N[i])
    return res

```

```

%timeit apply_integrate_f_cython_step3(df['a'].to_numpy(), df['b'].to_numpy(), df['N'])
# 29.2 ms ± 152 µs per loop (mean ± std. dev. of 7 runs, 10 loops each)

```

## Cython: Adding data type annotation to local variables and return (step 4)

Last step, we can add type annotation to local variables within functions and the return value.

[Skip to content](#)

Cython

```

%%cython -a

import cython
import numpy as np

@cython.cfunc
def f_cython_step4(x: cython.double) -> cython.double:
    return x ** 2 - x

@cython.cfunc
def integrate_f_cython_step4(
    a: cython.double,
    b: cython.double,
    N: cython.long
) -> cython.double:
    s: cython.double
    dx: cython.double
    i: cython.long

    s = 0
    dx = (b - a) / N

    for i in range(N):
        s += f_cython_step4(a + i * dx)
    return s * dx

@cython.ccall
def apply_integrate_f_cython_step4(
    col_a: cython.double[:],
    col_b: cython.double[:],
    col_N: cython.long[:]
) -> cython.double[:]:
    n: cython.int
    i: cython.int
    res: cython.double[:]

    n = len(col_N)
    res = np.empty(n, dtype=np.float64)
    for i in range(n):
        res[i] = integrate_f_cython_step4(col_a[i], col_b[i], col_N[i])
    return res

```

[Skip to content](#)

```

In [1]: %timeit integrate_f_cython_step4(df['a'].to_numpy(), df['b'].to_numpy(), df['N'])
# 471 µs ± 7.38 µs per loop (mean ± std. dev. of 7 runs, 1,000 loops each)

```

Now it is ~200 times faster than the baseline Python implementation, and all we have done is to add type declarations on the Python code!

Yellow lines hint at Python interaction.

Click on a line that starts with a "+" to see the C code that Cython generated for it.

```
01:  
+02: import cython  
+03: import numpy as np  
04:  
+05: @cython.cfunc  
06: def f_cython_step4(x: cython.double) -> cython.double:  
+07:     return x ** 2 - x  
08:  
+09: @cython.cfunc  
10: def integrate_f_cython_step4(  
11:     a: cython.double,  
12:     b: cython.double,  
13:     N: cython.long  
14: ) -> cython.double:  
15:     s: cython.double  
16:     dx: cython.double  
17:     i: cython.long  
18:  
+19:     s = 0  
+20:     dx = (b - a) / N  
21:  
+22:     for i in range(N):  
+23:         s += f_cython_step4(a + i * dx)  
+24:     return s * dx  
25:  
+26: @cython.ccall  
27: def apply_integrate_f_cython_step4(  
28:     col_a: cython.double[:],  
29:     col_b: cython.double[:],  
30:     col_N: cython.long[:]  
31: ) -> cython.double[:]:  
32:     n: cython.int  
33:     i: cython.int  
34:     res: cython.double[:]  
35:  
+36:     n = len(col_N)  
+37:     res = np.empty(n, dtype=np.float64)  
+38:     for i in range(n):  
+39:         res[i] = integrate_f_cython_step4(col_a[i], col_b[i], col_N[i])  
+40:     return res
```

WE INDEED SEE MUCH LESS PYTHON INTERACTION IN THE CODE FROM STEP 1 TO STEP 4.

## Other useful features

There are some useful (and possibly advanced) features which are not covered in this episode. Some of these features are called [magic attributes](#). Here are a few:

- `cython.cimports` package for importing and calling C libraries such as [libc.math](#).

[Skip to content](#)

## Note

Differences between `import` (for Python) and `cimport` (for Cython) statements

- `import` gives access to Python libraries, functions or attributes
- `cimport` gives access to C libraries, functions or attributes

In case of Numpy it is common to use the following, and Cython will internally handle this ambiguity.

Pure Python

Cython

```
from cython.cimports.libc.stdlib import malloc, free # Allocate and free memory
from cython.cimports.libc import math # For math functions like sin, cos etc.
from cython.cimports import numpy as np # access to NumPy C API
```

- `cython.nogil`, which can act both as a decorator or context-manager, to manage the GIL (Global Interpreter Lock). See [Cython and the GIL](#).
- `@cython.boundscheck(False)` and `@cython.wraparound(False)` decorators to tune indexing of Numpy array. See [Cython for NumPy users](#).
- `@cython.cclass` to declare [Extension Types](#) which behave similar to Python classes.

In addition to the above Cython can also,

- [augment with .pxd files](#) where the Python code is kept as it is and the `.pxd` file describes the type annotation. In this form `.pxd` is very similar in function to a C/C++ header file or `.pyi` Python type annotation file,
- create parallel code using [parallel blocks](#) and `prange` iterator for element-wise parallel operation or reductions based on OpenMP threads (see [Writing parallel code with Cython](#)).

[Skip to content](#)

## | i Demo

Here is a code which showcases most of the features above, except the `@cython.cclass` feature and the use of `.pxd` files.

Pure Python

Cython

Numpy

Naive Python implementation

```
import cython
from cython.parallel import parallel, prange
from cython.cimports.libc.math import sqrt

@cython.boundscheck(False)
@cython.wraparound(False)
def normalize(x: cython.double[:]):
    """Normalize a 1D array by dividing all its elements using its root-mean-square (RMS) value."""
    i: cython.Py_ssize_t
    total: cython.double = 0
    norm: cython.double
    with cython.nogil, parallel():
        for i in prange(x.shape[0]):
            total += x[i]*x[i]
    norm = sqrt(total)
    for i in prange(x.shape[0]):
        x[i] /= norm
```

## | i Note

If you compare performance of the the Cython code versus the Numpy code, you might observe that it is either on-par, or slightly worse than Numpy. This is because Numpy vectorized operations also makes use of OpenMP parallelism and is heavily optimized. Nevertheless, it is orders of magnitude better than a naive implementation.

## Conclusions

### | Keypoints

- Cython is a versatile, general purpose compiler for Python code
- Cython is a great way to write high-performance code in Python where algorithms are not available in scientific libraries like Numpy and Scipy and require custom implementation

### | i See also

In order to make Cython code reusable often some packaging is necessary. The compilation to binary extension can either happen during the packaging itself, or during installation of a Python package. To learn more about how to skip to content

sions, read the following guides:

- *cython* Python packaging guide's page on [build tools](#)
- *Python packaging user guide*'s page on [packaging binary extensions](#)

# Quick Reference

## Instructor's guide

**Why we teach this lesson**

**Intended learning outcomes**

**Timing**

**Preparing exercises**

e.g. what to do the day before to set up common repositories.

**Other practical aspects**

**Interesting questions you might get**

**Typical pitfalls**

## Learning outcomes

**FIXME**

This material is for ...

By the end of this module, learners should:

- ...
- ...

## See also



Credit

**FIXME**

Don't forget to check out additional course materials from ...

[Skip to content](#)

## License

### CC BY-SA for media and pedagogical material

Copyright © 2025 XXX. This material is released by XXX under the Creative Commons Attribution-ShareAlike 4.0 International (CC BY-SA 4.0).

**Canonical URL:** <https://creativecommons.org/licenses/by-sa/4.0/>

[See the legal code](#)

### You are free to

1. **Share** — copy and redistribute the material in any medium or format for any purpose, even commercially.
2. **Adapt** — remix, transform, and build upon the material for any purpose, even commercially.
3. The licensor cannot revoke these freedoms as long as you follow the license terms.

### Under the following terms

1. **Attribution** — You must give [appropriate credit](#), provide a link to the license, and [indicate if changes were made](#). You may do so in any reasonable manner, but not in any way that suggests the licensor endorses you or your use.
2. **ShareAlike** — If you remix, transform, or build upon the material, you must distribute your contributions under the [same license](#) as the original.
3. **No additional restrictions** — You may not apply legal terms or [technological measures](#) that legally restrict others from doing anything the license permits.

### Notices

You do not have to comply with the license for elements of the material in the public domain or where your use is permitted by an applicable [exception or limitation](#).

No warranties are given. The license may not give you all of the permissions necessary for your intended use. For example, other rights such as [publicity, privacy, or moral rights](#) may limit how you use the material.

This deed highlights only some of the key features and terms of the actual license. It is not a license and has no legal value. You should carefully review all of the terms and conditions of the actual license before using the licensed material.

[Skip to content](#)

## MIT for source code and code snippets

MIT License

Copyright (c) 2025, ENCCS project, The contributors

Permission is hereby granted, free of charge, to any person obtaining a copy of this software and associated documentation files (the "Software"), to deal in the Software without restriction, including without limitation the rights to use, copy, modify, merge, publish, distribute, sublicense, and/or sell copies of the Software, and to permit persons to whom the Software is furnished to do so, subject to the following conditions:

The above copyright notice and this permission notice shall be included in all copies or substantial portions of the Software.

THE SOFTWARE IS PROVIDED "AS IS", WITHOUT WARRANTY OF ANY KIND, EXPRESS OR IMPLIED, INCLUDING BUT NOT LIMITED TO THE WARRANTIES OF MERCHANTABILITY, FITNESS FOR A PARTICULAR PURPOSE AND NONINFRINGEMENT. IN NO EVENT SHALL THE AUTHORS OR COPYRIGHT HOLDERS BE LIABLE FOR ANY CLAIM, DAMAGES OR OTHER LIABILITY, WHETHER IN AN ACTION OF CONTRACT, TORT OR OTHERWISE, ARISING FROM, OUT OF OR IN CONNECTION WITH THE SOFTWARE OR THE USE OR OTHER DEALINGS IN THE SOFTWARE.



Copyright © 2025, ENCCS, The contributors  
Made with [Sphinx](#) and @pradyunsg's [Furo](#)